

Machine Learning

12 – Temporal Probability Models

SS 2018

Gunther Heidemann

Temporal probability models

- Modeling uncertainty for temporal processes
- Markov processes
- Temporal inference:
 - Filtering
 - Prediction
 - Smoothing
 - Most likely explanation (Viterbi algorithm)
- Brief overview of dynamic Bayes Networks (DBN)

Textbook:

Stuart Russell, Peter Norvig: *Artificial Intelligence*, Pearson

Aim: Track and predict processes over time

Examples: Diabetes management; motor management

Description of a process:

- Description by discrete **states** (using discrete time t).
- State at t is specified by a set X_t of **unobservable variables** (= this is the problem!), e.g.,

$$X_t = \{BloodSugar_t, StomachContents_t\}.$$

- The state becomes visible only by a set E_t of observable evidence variables, e.g.,

$$E_t = \{MeasuredBloodSugar_t, FoodEaten_t, PulseRate_t\}.$$

Notation for a span of time:

$$X_{a:b} = X_a, X_{a+1}, \dots, X_{b-1}, X_b.$$

Markov assumption for a process described by variable X_t (discrete time):

X_t depends only on a **bounded subset** of the variables $X_{0:t-1}$.

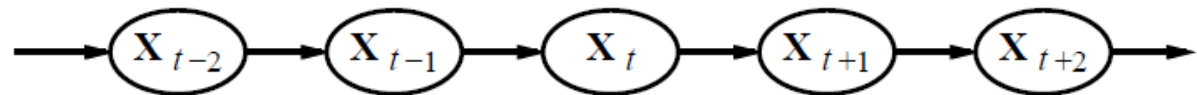
First order Markov process:

$$P(X_t | X_{0:t-1}) = P(X_t | X_{t-1})$$

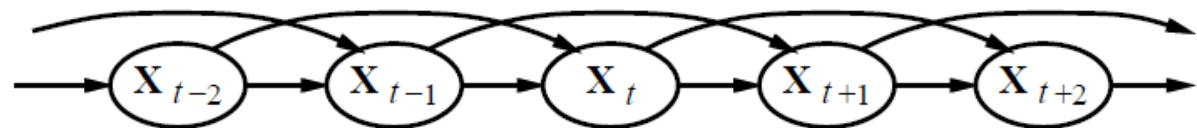
Second order Markov process:

$$P(X_t | X_{0:t-1}) = P(X_t | X_{t-2}, X_{t-1})$$

First order



Second order



[RN]

Real world:

- Assumption of a first order Markov process simplifies modelling,
- but does usually not strictly hold.

Improvements:

1. Assume Markov process of higher order.
2. Augment knowledge about the state by additional evidence variables.

Example: Moving robot

Augment state description $(position_t, velocity_t)$ by $battery_t$.

Markov assumption for sensors:

Sensor output depends only on the current value of the evidence variable observed by the sensor.

$$P(E_t | X_{0:t}, E_{0:t-1}) = P(E_t | X_t)$$

Example: Speedometer

Observed velocity E_t depends only on current velocity X_t , not on past values X_{t-n} .

Counter example: Water gauge for hydroponics.

- Unobservable variable: Height of water X_t
- Observed evidence variable: Measured height E_t
- If X exceeds a maximum height for weeks, algae block the water gauge → no Markov sensor.

Stationary process:

The world is changing, but the rules underlying both this change and its observation remain the same.

$$\rightarrow P(X_t | X_{0:t-1}, t) = P(X_t | X_{0:t-1}).$$

Together with the Markov assumptions for the process and the sensor, all we need to describe the measurements are the

transition model $P(X_t | X_{t-1})$ and the

sensor model (model of observation) $P(E_t | X_t)$,

which both remain fixed for all t .

Example of a stationary state transition model:

- X is water level of hydroponics.
- Add approx. 1l water every week (Gaussian distribution with mean = 1l).

Counter example:

- In summer, the weather becomes hot and the average is increased to 1,2l (rules of the world have changed).

Example for a stationary sensor model:

- Weight measurement using spring:
- X is the weight, E the weight measurement (which may exhibit a systematic but stationary error).

Counter example:

- Over time, the spring loses strength.

Problem:

- Stationary first order Markov process with known **state transition model** $\mathbf{P}(X_t | X_{t-1})$.
- Markov sensor with known **sensor model** $\mathbf{P}(E_t | X_t)$.

Inference:

Several types of inference which are different mixtures of two basic problems:

- Infer a past or the current value of an unobservable state variable X from the observable evidence variables.
- Prediction.

The type of inference depends on the times for which the probability distribution of X is inferred from past to current evidence values e .

Let T denote the current time (“now”):

1. **Filtering:** $P(X_T | e_{1:T})$

Infer probability distribution P of current state X_T from current evidence value e_T and past evidence values $e_{1:T-1}$ (but note $x_1 \dots x_{T-1}$ are unknown).

$P(X_T | e_{1:T})$ is called a *belief state*. This is the input for the decision process of a rational agent.

2. **Prediction:** $P(X_{T+K} | e_{1:T}), K > 0$

Predict future value of X_{T+K} from evidence values $e_{1:T}$. Same as filtering, but the future evidence values $e_{T+1:T+K}$ are unknown.

3. Smoothing: $\mathbf{P}(X_K \mid e_{1:T})$, $0 < K < T$

Infer the probability distribution of X at the past time K from earlier evidence values $e_{1:K-1}$ and later evidence values $e_{K+1:T}$.

Yields better estimate for past states than filtering.

4. Most likely explanation: $\arg \max_{X_{1:T}} \mathbf{P}(X_{1:T} \mid e_{1:T})$

Infer all $X_{1:T}$ from all evidence values $e_{1:T}$.

Example: Speech recognition.

Given: **Transition model** $\mathbf{P}(X_t | X_{t-1})$ and **sensor model** $\mathbf{P}(E_t | X_t)$.

Wanted: Probabilities for X_T from $e_{1:T}$ (better than mere estimation of X_T from e_T).

Principle: **Recursive state estimation** algorithm which starts with an assumption for X_0 and can infer values at $t+1$ from t .

For $t = 0, 1, \dots, T-1$:

$$\begin{aligned}
 & \mathbf{P}(X_{t+1} | e_{1:t+1}) \\
 &= \mathbf{P}(X_{t+1} | e_{1:t}, e_{t+1}) \\
 &= \alpha \mathbf{P}(e_{t+1} | X_{t+1}, e_{1:t}) \mathbf{P}(X_{t+1} | e_{1:t}) && \text{(Bayes)} \\
 &= \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \mathbf{P}(X_{t+1} | e_{1:t}) && \text{(Markov sensor)} \\
 &= \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_{x_t} \mathbf{P}(X_{t+1} | x_t, e_{1:t}) \mathbf{P}(x_t | e_{1:t}) && \text{(summing out } X_t) \\
 &= \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_{x_t} \mathbf{P}(X_{t+1} | x_t) \mathbf{P}(x_t | e_{1:t}) && \text{(Markov process)} \\
 &= \alpha \text{ sensor model}_{t+1} \sum_{x_t} \text{transition model}_{t+1,t} \text{probability distribution}_t
 \end{aligned}$$

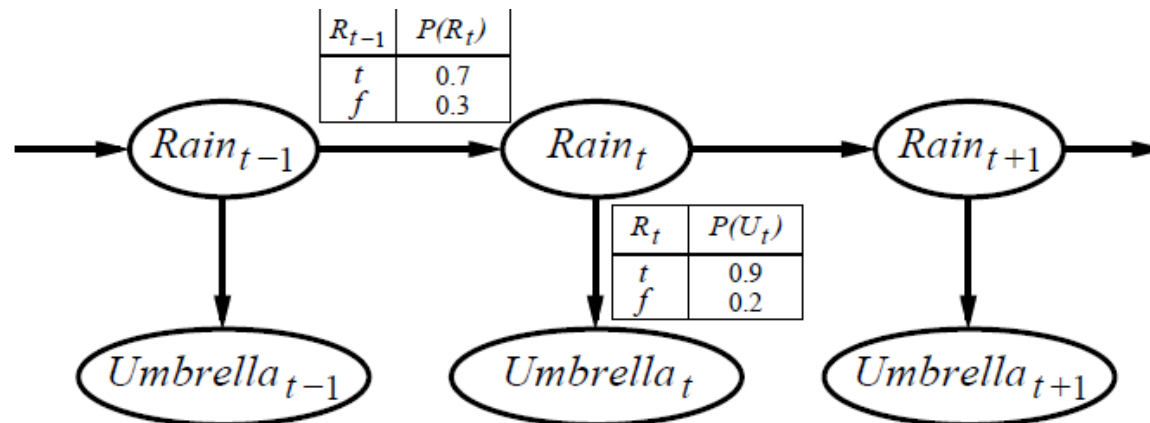
Structure of recursion:

$$f_{1:t+1} = \text{Forward}(f_{1:t}, e_{t+1}) \quad \text{with} \quad f_{1:t} = P(X_t | e_{1:t})$$

Requires **constant** time and memory for each step, independent of t .

Example [RN]:

- Living in a bunker, time steps are days.
- Only your boss is allowed to go outside.
- You can infer whether it is raining (X_t) only from his umbrella.



[RN]

Transition model:

$$\mathbf{P}(X_t | X_{t-1}) = \mathbf{P}(\text{Rain}_t | \text{Rain}_{t-1}) \quad \text{with} \quad \begin{aligned} \mathbf{P}(X_t | X_{t-1} = \text{true}) &= \langle 0.7, 0.3 \rangle \\ \mathbf{P}(X_t | X_{t-1} = \text{false}) &= \langle 0.3, 0.7 \rangle \end{aligned}$$

Sensor model:

$$\mathbf{P}(E_t | X_t) = \mathbf{P}(\text{Umbrella}_t | \text{Rain}_t) \quad \text{with} \quad \begin{aligned} \mathbf{P}(E_t | X_t = \text{true}) &= \langle 0.9, 0.1 \rangle \\ \mathbf{P}(E_t | X_t = \text{false}) &= \langle 0.2, 0.8 \rangle \end{aligned}$$

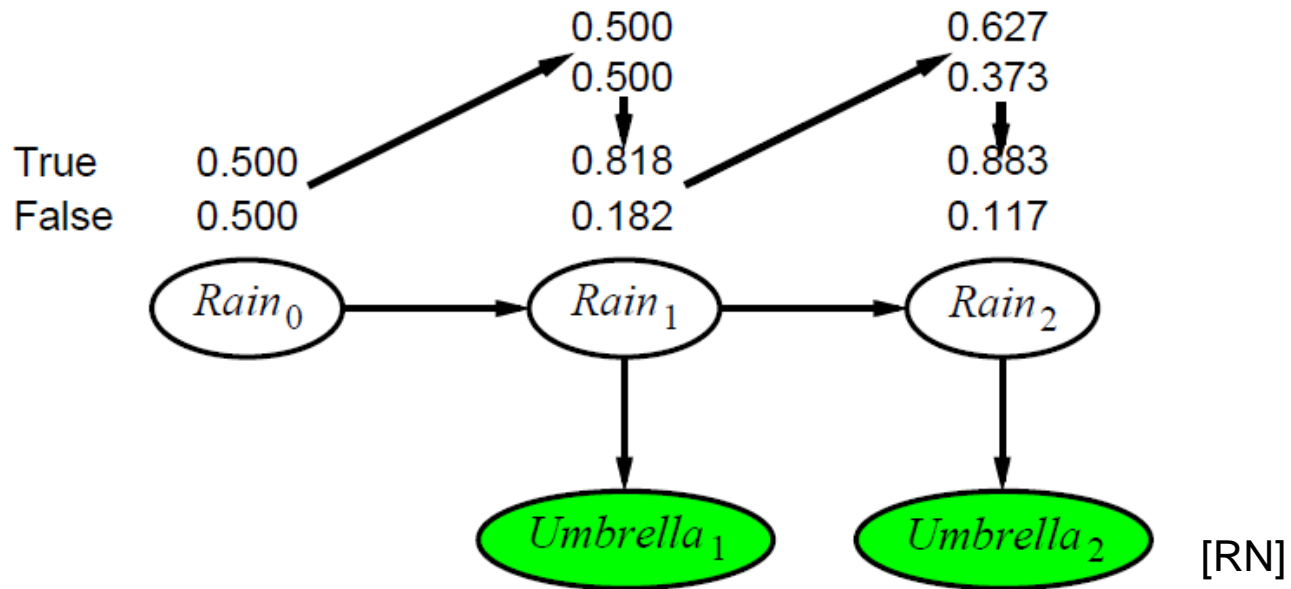
What is the probability for rain $\mathbf{P}(X_2)$ on the second day ($T=2$), if

- Day 1: Umbrella $e_1 = \text{true}$.
- Day 2: Umbrella $e_2 = \text{true}$.
- We need an assumption about the probability of rain for day 0:

$$\mathbf{P}(X_0) = \langle 0.5, 0.5 \rangle$$

$$\begin{aligned}
 \mathbf{P}(X_{t+1} | e_{1:t+1}) &= \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_{x_t} \mathbf{P}(X_{t+1} | x_t) \mathbf{P}(x_t | e_{1:t}) \\
 \mathbf{P}(X_2 | e_{1:2}) &= \alpha \mathbf{P}(e_2 | X_2) \sum_{x_1} \mathbf{P}(X_2 | x_1) \mathbf{P}(x_1 | e_1) \\
 &= \alpha \mathbf{P}(e_2 | X_2) [\mathbf{P}(X_2 | x_1=t) 0.818 + \mathbf{P}(X_2 | x_1=f) 0.182] \\
 &= \alpha \langle 0.9, 0.2 \rangle [\langle 0.7, 0.3 \rangle 0.818 + \langle 0.3, 0.7 \rangle 0.182] \\
 &= \alpha \langle 0.9, 0.2 \rangle \langle 0.627, 0.373 \rangle = \alpha \langle 0.565, 0.075 \rangle \\
 &= \langle 0.883, 0.117 \rangle
 \end{aligned}$$

$$\begin{aligned}
 \mathbf{P}(X_1 | e_1) &= \alpha' \mathbf{P}(e_1 | X_1) \sum_{x_0} \mathbf{P}(X_1 | x_0) \mathbf{P}(x_0) \\
 &= \alpha' \mathbf{P}(e_1 | X_1) [\mathbf{P}(X_1 | x_0=t) \mathbf{P}(x_0=t) + \mathbf{P}(X_1 | x_0=f) \mathbf{P}(x_0=f)] \\
 &= \alpha' \mathbf{P}(e_1 | X_1) [\langle 0.7, 0.3 \rangle 0.5 + \langle 0.3, 0.7 \rangle 0.5] \\
 &= \alpha' \mathbf{P}(e_1 | X_1) \langle 0.5, 0.5 \rangle \\
 &= \alpha' \langle 0.9, 0.2 \rangle \langle 0.5, 0.5 \rangle = \alpha' \langle 0.45, 0.1 \rangle = \langle 0.818, 0.182 \rangle
 \end{aligned}$$



Given: Transition model $\mathbf{P}(X_t | X_{t-1})$ and sensor model $\mathbf{P}(E_t | X_t)$.

Wanted: X_{T+K} from $e_{1:T}$.

Idea:

- Filtering up to T
- After T : K further steps without new evidences (we lack $e_{T+1:T+K}$).
- New recursion: $T+k \rightarrow T+k+1$

For $k = 0, 1, \dots, K-1$:

$$\begin{aligned} \mathbf{P}(X_{T+k+1} | e_{1:T}) &= \sum_{x_{T+k}} \mathbf{P}(X_{T+k+1} | x_{T+k}) \mathbf{P}(x_{T+k} | e_{1:T}) \\ &= \sum_{x_{T+k}} \text{transition model}_{T+k+1, T+k} \text{distribution}_{T+k} \end{aligned}$$

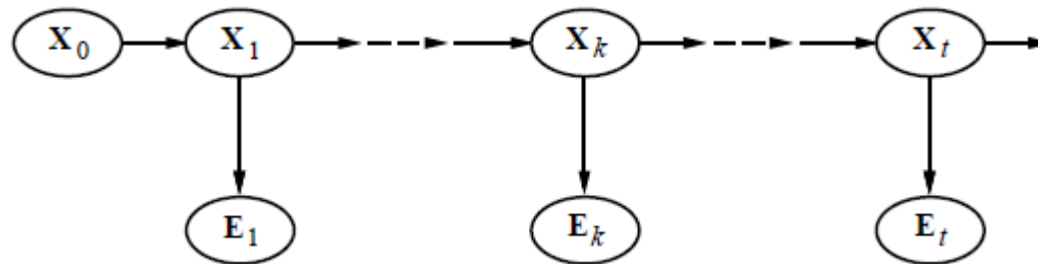
The farther we predict the future without new evidences, the more the distribution is dominated by the transition model.

Given: Transition model $P(X_t | X_{t-1})$ and sensor model $P(E_t | X_t)$.

Wanted: X_K from $e_{1:T}$ with $1 \leq K < T$.

Idea: „Forward-Backward-Algorithm“:

Filtering from 1 to K and „backward filtering“ from T to K .



[RN]

Procedure:

1. Split problem into forward- and backward-part.
2. Forward filtering: Already known.
3. Backward filtering: New.

Use evidences $e_{1:K}$ up to K and later ones $e_{K+1:T}$ separately:

$$\begin{aligned}
 P(X_K | e_{1:T}) &= P(X_K | e_{1:K}, e_{K+1:T}) \\
 &= \alpha P(X_K | e_{1:K}) P(e_{K+1:T} | X_K, e_{1:K}) && \text{(Bayes)} \\
 &= \alpha P(X_K | e_{1:K}) P(e_{K+1:T} | X_K) && \text{(Markov sensor)} \\
 &= \alpha f_{1:K} b_{K+1:T}
 \end{aligned}$$

Backward recursion for $k = T-1, T-2, \dots, K+1, K$:

$$\begin{aligned}
 P(e_{k+1:T} | X_k) &= \sum_{x_{k+1}} P(e_{k+1:T} | X_k, x_{k+1}) P(x_{k+1} | X_k) \\
 &= \sum_{x_{k+1}} P(e_{k+1:T} | x_{k+1}) P(x_{k+1} | X_k) && \text{(cond. ind.)} \\
 &= \sum_{x_{k+1}} P(e_{k+1}, e_{k+2:T} | x_{k+1}) P(x_{k+1} | X_k) \\
 &= \sum_{x_{k+1}} P(e_{k+1} | x_{k+1}) P(e_{k+2:T} | x_{k+1}) P(x_{k+1} | X_k)
 \end{aligned}$$

Structure: $b_{k+1:T} = \text{Backward}(b_{k+2:T}, e_{k+1:T})$ with $b_{k+1:T} = P(e_{k+1:T} | X_k)$

Rain – umbrella domain as before:

$$\mathbf{P}(X_t | X_{t-1} = \text{true}) = \langle 0.7, 0.3 \rangle \quad \mathbf{P}(E_t | X_t = \text{true}) = \langle 0.9, 0.1 \rangle$$

$$\mathbf{P}(X_t | X_{t-1} = \text{false}) = \langle 0.3, 0.7 \rangle \quad \mathbf{P}(E_t | X_t = \text{false}) = \langle 0.2, 0.8 \rangle$$

$$\mathbf{P}(X_0) = \langle 0.5, 0.5 \rangle; e_1 = \text{true}, e_2 = \text{true}. \quad \mathbf{P}(X_1 | e_{1:2}) = ?$$

In general: $\mathbf{P}(X_K | e_{1:T}) = \alpha \mathbf{P}(X_K | e_{1:K}) \mathbf{P}(e_{K+1:T} | X_K)$, here: $T=2, K=1$

Here: $\mathbf{P}(X_1 | e_{1:2}) = \alpha \mathbf{P}(X_1 | e_1) \mathbf{P}(e_2 | X_1)$

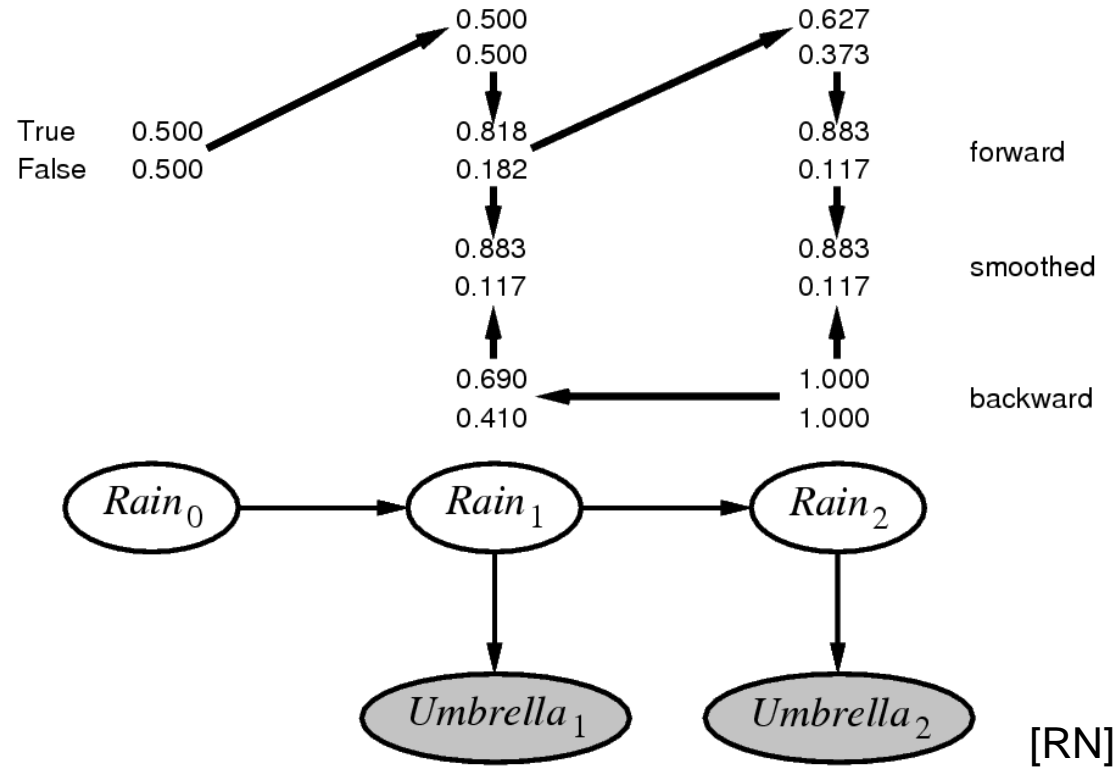
By filtering: $\mathbf{P}(X_1 | e_1) = \langle 0.818, 0.182 \rangle$

$$\mathbf{P}(e_{k+1:T} | X_k) = \sum_{x_{k+1}} \mathbf{P}(e_{k+1} | x_{k+1}) \mathbf{P}(e_{k+2:T} | x_{k+1}) \mathbf{P}(x_{k+1} | X_k)$$

$$\begin{aligned} \mathbf{P}(e_2 | X_1) &= \sum_{x_2} \mathbf{P}(e_2 | x_2) \mathbf{P}(e_{3:2} | x_2) \mathbf{P}(x_2 | X_1) \text{ with } \mathbf{P}(e_{3:2} | x_2) = 1. \\ &= 0.9 \cdot 1 \cdot \langle 0.7, 0.3 \rangle + 0.2 \cdot 1 \cdot \langle 0.3, 0.7 \rangle = \langle 0.69, 0.41 \rangle \end{aligned}$$

$$\mathbf{P}(X_1 | e_{1:2}) = \alpha \langle 0.818, 0.182 \rangle \cdot \langle 0.69, 0.41 \rangle = \langle 0.883, 0.117 \rangle$$

Smoothing



- Problem: Find the most likely explanation for a sequence of observed events.
- More precisely: Find the most likely sequence of hidden states that would cause the observed sequence of evidences.
- Example: For a boolean variable, for T steps there are 2^T possible sequences of states.
- Naive approach: Compute for each state in separation the probabilities using smoothing.
- **But: The most likely sequence \neq the sequence of most likely states!**
- The most likely sequence requires maximizing the joint probability (not the isolated probabilities) !
- Solution: *Viterbi algorithm*.
- Applications: Cell phones, WLAN, hard disks, speech recognition.

Most likely path to x_{t+1} = most likely path to x_t plus another step:

$$\max_{x_1 \dots x_t} \mathbf{P}(x_1, \dots, x_t, X_{t+1} \mid e_{1:t+1}) = \\ \alpha \mathbf{P}(e_{t+1} \mid X_{t+1}) \max_{x_t} [\mathbf{P}(X_{t+1} \mid x_t) \max_{x_1 \dots x_{t-1}} \mathbf{P}(x_1, \dots, x_{t-1}, x_t \mid e_{1:t})]$$

Like filtering ($f_{1:t+1} = \alpha \mathbf{P}(e_{t+1} \mid X_{t+1}) \sum_{x_t} \mathbf{P}(X_{t+1} \mid x_t) f_{1:t}$),

but:

1. $f_{1:t} = \mathbf{P}(X_t \mid e_{1:t})$ is replaced by

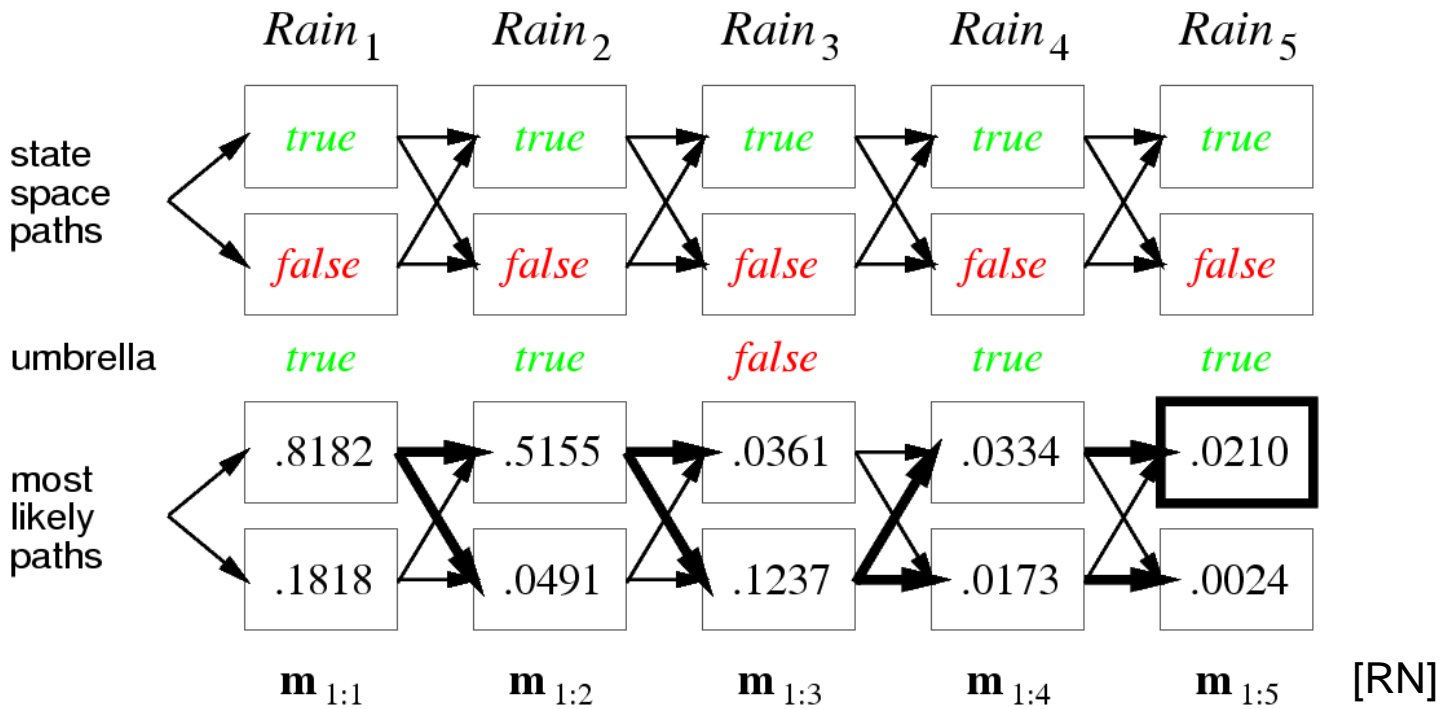
$$m_{1:t} = \max_{x_1 \dots x_{t-1}} \mathbf{P}(x_1, \dots, x_{t-1}, X_t \mid e_{1:t}),$$

i.e. $m_{1:t}(i)$ is the probability of the most likely path to state i .

2. Replace sum over x_t by maximizing over x_t (Viterbi algorithm):

$$m_{1:t+1} = \mathbf{P}(e_{t+1} \mid X_{t+1}) \max_{x_t} [\mathbf{P}(X_{t+1} \mid x_t) m_{1:t}]$$

1. Compute all $m_{1:t}$ successively. For each state, memorize the best previous state (thick arrows).
2. Choose the most likely state for time t .
3. Go back to the best previous state and so on.



- So far: Transition model and sensor model were given by the experiment, no formal description.
- If a Markov process is described by states with a single variable:
Hidden-Markov-Modell (HMM)
- Modeling temporal process and its evidences by two random processes:
 1. Random process: Markov chain with one hidden variable the transitions of which are described by probabilities.
 2. Random process: A Markov sensor provides evidences of the hidden variable.

Definition of a HMM:

- Let X_t be a **single** discrete random variable taking values (states) $\{s_1 \dots s_n\}$, and
- E_t its evidence variable with values (possible observations) $\{e_1 \dots e_m\}$. The **matrix** $T_{ij} = P(X_t = s_j \mid X_{t-1} = s_i)$ describes the probabilities for **state transitions**.
- The matrix $O_{ij} = P(e_j \mid s_i)$ is the **observation matrix** of probabilities that the Markov sensor yields observation e_j for state s_i .
- Starting distribution for X_0 .

A HMM is stationary if T and O do not change over time.

With the

transition matrix $T_{ij} = P(X_t = j \mid X_{t-1} = i)$ and the

observation matrix $(O_t)_{ij} = P(e_t \mid X_t = i)$

we can simplify, e.g., smoothing using matrix notation:

$$f_{1:t+1} = \alpha O_{t+1} T^T f_{1:t}$$

in place of

$$\mathbf{P}(X_{t+1} \mid e_{1:t+1}) = \alpha \mathbf{P}(e_{t+1} \mid X_{t+1}) \sum_{x_t} \mathbf{P}(X_{t+1} \mid x_t) \mathbf{P}(x_t \mid e_{1:t})$$

and

$$b_{k+1:t} = T O_{k+1} b_{k+2:t}$$

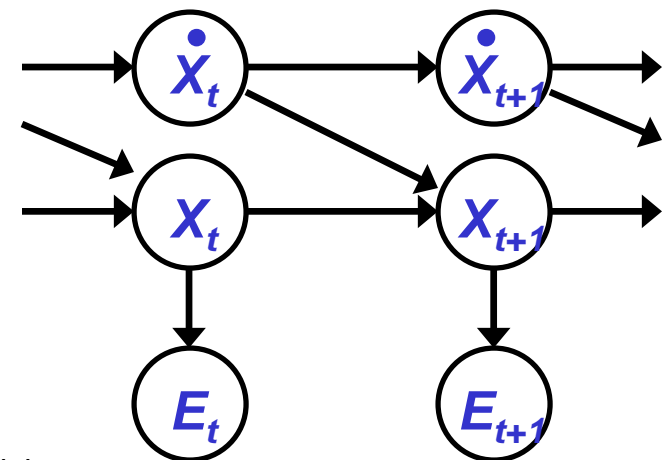
in place of

$$\mathbf{P}(e_{k+1:T} \mid X_k) = \sum_{x_{k+1}} \mathbf{P}(e_{k+1} \mid x_{k+1}) \mathbf{P}(e_{k+2:T} \mid x_{k+1}) \mathbf{P}(x_{k+1} \mid X_k).$$

- So far: No continuous variables.
- Kalman filtering provides a model for systems with continuous variables, in particular, time dependent variables.
- Example: Trajectory tracking. Position and its temporal derivative (velocity) are considered random variables.

A bird is flying through a forest. Try to predict its trajectory though it is partially hidden behind trees.

- Other examples: Planets, robots, ecosystems, markets, fusion GPS – inertial sensors.
- Bayes network for linear dynamical system with position X_t and position measurement E_t :



Example: 1D trajectory

- Observe X -coordinate
- Observation at intervals Δt .
- Assumption: Velocity is approximately constant.

Simple trajectory prediction: $X_{t+\Delta t} = X_t + \dot{X} \Delta t$.

To account for measurement errors and non-constant velocity we assume an error with Gaussian distribution:

$$P(X_{t+\Delta t} = x_{t+\Delta t} \mid X_t = x_t, \dot{X}_t = \dot{x}_t) = N(x_t + \dot{x}_t \Delta t, \sigma, x_{t+\Delta t})$$

with $N(x_0, \sigma, x) = \alpha \exp(-\frac{1}{2} (x - x_0)^2 / \sigma)$.

Assumptions:

- Gaussian a-priori distribution
- Linear Gaussian transition model
- Linear Gaussian observation model.

Prediction:

If $\mathbf{P}(X_t | e_{1:t})$ has a Gaussian distribution, then the predicted distribution is also Gaussian:

$$\mathbf{P}(X_{t+1} | e_{1:t}) = \int_{X_t} \mathbf{P}(X_{t+1} | x_t) \mathbf{P}(x_t | e_{1:t}) dx_t$$

With $\mathbf{P}(X_{t+1} | e_{1:t})$ also we also have a Gaussian for

$$\mathbf{P}(X_{t+1} | e_{1:t+1}) = \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \mathbf{P}(X_{t+1} | e_{1:t}).$$

Hence, $\mathbf{P}(X_t | e_{1:t})$ is a multivariate Gaussian $N(\mu_t, \Sigma_t)$ for all t with mean μ and covariance matrix Σ .

- $P(X_t | e_{1:t})$ is (and stays!) Gaussian. Its parameters (mean, covariance) change over time.
- Thus $P(X_t | e_{1:t})$ can be described with the **same number of parameters** for all times t .
- As the Gaussian may become arbitrarily broad, the usable information on X may become very small, but ...
- ... at least, this small amount of usable information is still encoded in the same number of parameters.
- For the **general** case (non-linear, non-Gaussian) this does not hold: In general, the effort for the description of the posterior **grows** over time!

Kalman filtering: 1D random walk

- Gaussian random walk along X -axis, X_t is the random variable.
- Prior distribution (initial position measured with limited accuracy):

$$P(x_0) = N(\mu_0, \sigma_0, x_0).$$

- Transition model (walk along random path):

$$P(x_{t+1} | x_t) = N(x_t, \sigma_x, x_{t+1}).$$

- Observation model (position measurement with limited accuracy):

$$P(e_t | x_t) = N(e_t, \sigma_e, x_t)$$

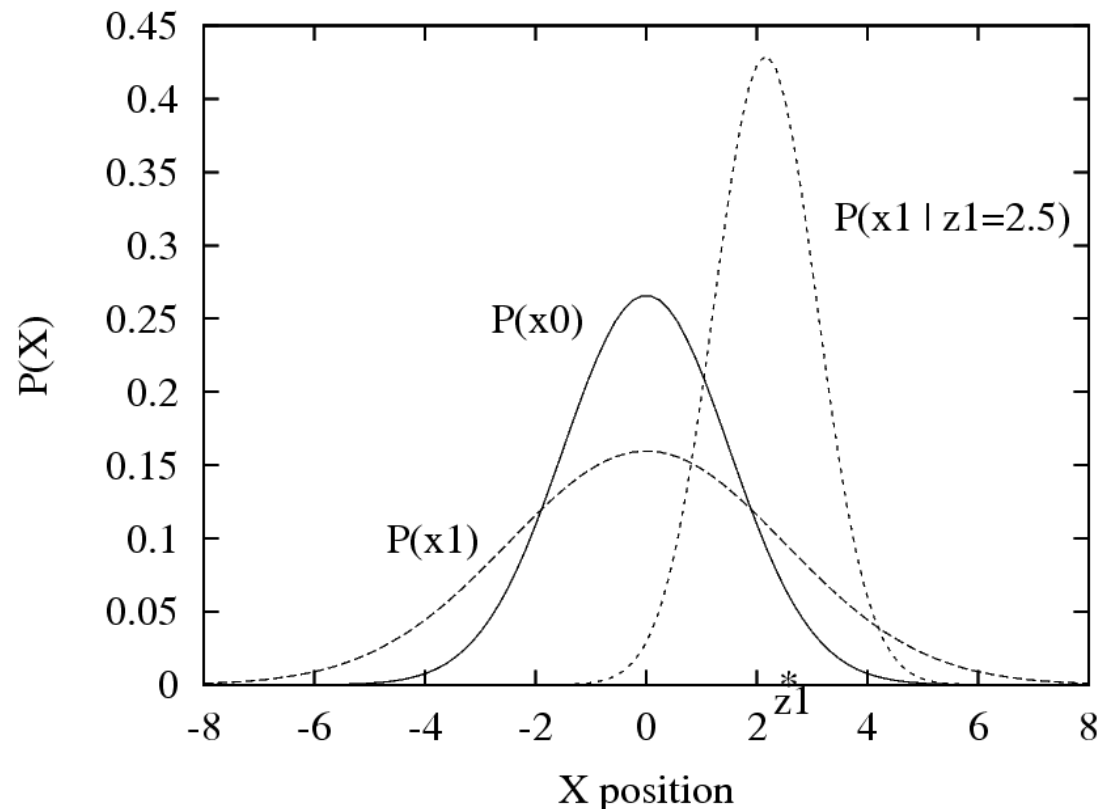
- First observation: e_1

$$P(x_1 | e_1) = N(\mu_1, \sigma_1, x_1) \quad \text{with} \quad \mu_1 = \frac{(\sigma_0^2 + \sigma_x^2)e_1 + \sigma_e^2\mu_0}{\sigma_0^2 + \sigma_x^2 + \sigma_e^2}, \quad \sigma_1^2 = \frac{(\sigma_0^2 + \sigma_x^2)\sigma_e^2}{\sigma_0^2 + \sigma_x^2 + \sigma_e^2}$$

- In general: $\mu_{t+1} = \frac{(\sigma_t^2 + \sigma_x^2)e_{t+1} + \sigma_e^2\mu_t}{\sigma_t^2 + \sigma_x^2 + \sigma_e^2}, \quad \sigma_{t+1}^2 = \frac{(\sigma_t^2 + \sigma_x^2)\sigma_e^2}{\sigma_t^2 + \sigma_x^2 + \sigma_e^2}$

Kalman filtering: 1D random walk

- Initial distribution: $\mu_0 = 0, \sigma_0 = 1$
- Transition cause by noise with $\sigma_x = 2.$
- Sensor noise: $\sigma_e = 1.$
- First observation: $e_1 = 2.5.$
- Prediction $P(x_1)$ is more flat than $P(x_0)$ due to the noisy transition.
- The mean μ_1 of $P(x_1 | e_1)$ is smaller than 2.5, because the prediction $P(x_1)$ is accounted for.



Vector \vec{x} of n random variables.

Vector of n observation values: \vec{e} .

Transition model: $\mathbf{P}(\vec{x}_{t+1} | \vec{x}_t) = \mathbf{N}(F\vec{x}_t, \Sigma_x, \vec{x}_{t+1})$

Observation model: $\mathbf{P}(\vec{e}_t | \vec{x}_t) = \mathbf{N}(H\vec{x}_t, \Sigma_e, \vec{e}_t)$

F : $n \times n$ - matrix of the linear transition model

H : $n \times n$ - matrix of the linear observation model

Σ_x : $n \times n$ - covariance matrix of the transition noise

Σ_e : $n \times n$ - covariance matrix of the observation noise

Gaussian with n variables:

$$\mathbf{N}(\vec{\mu}, \Sigma, \vec{x}) = \alpha \exp(-\frac{1}{2} (\vec{x} - \vec{\mu})^T \Sigma^{-1} (\vec{x} - \vec{\mu}))$$

Updating rule:

$$\vec{\mu}_{t+1} = F \vec{\mu}_t + K_{t+1} (\vec{e}_{t+1} - H F \vec{\mu}_t)$$

$$\Sigma_{t+1} = (1 - K_{t+1}) L$$

with

$$L = F \Sigma_t F^T + \Sigma_x$$

$$K_{t+1} = L H^T (H L H^T + \Sigma_e)^{-1}$$

K is the **Kalman-Gain matrix**.

Interpretation:

$F \vec{\mu}_t$: Predicted $\vec{\mu}$
(according to linear model)

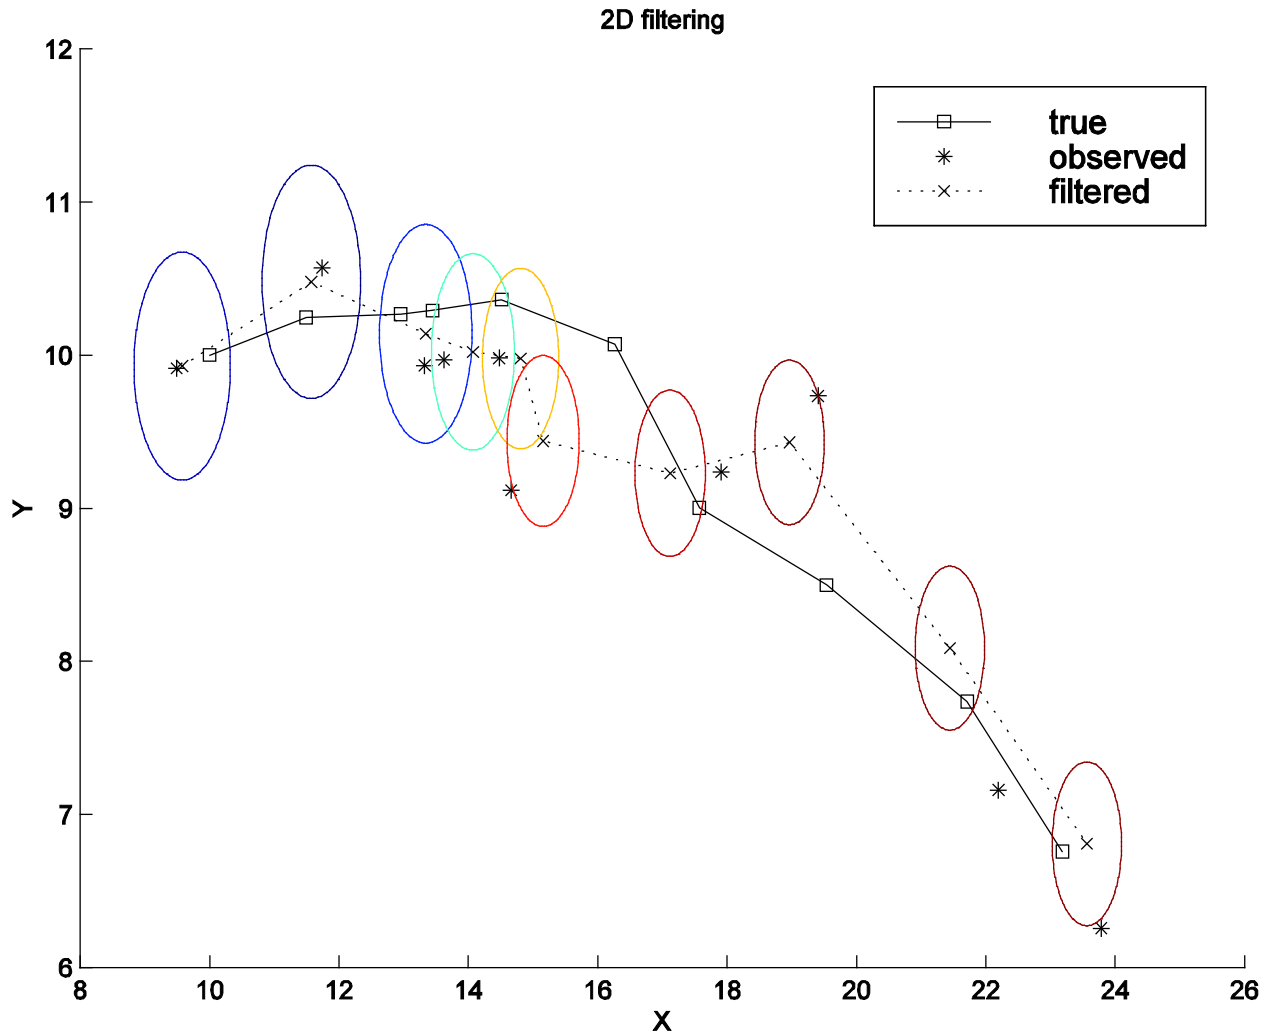
$H F \vec{\mu}_t$: Predicted observation

$\vec{e}_{t+1} - H F \vec{\mu}_t$: Difference between prediction and observation

K_{t+1} : Confidence we have in the observation, used as a weight for comparison with the linear prediction

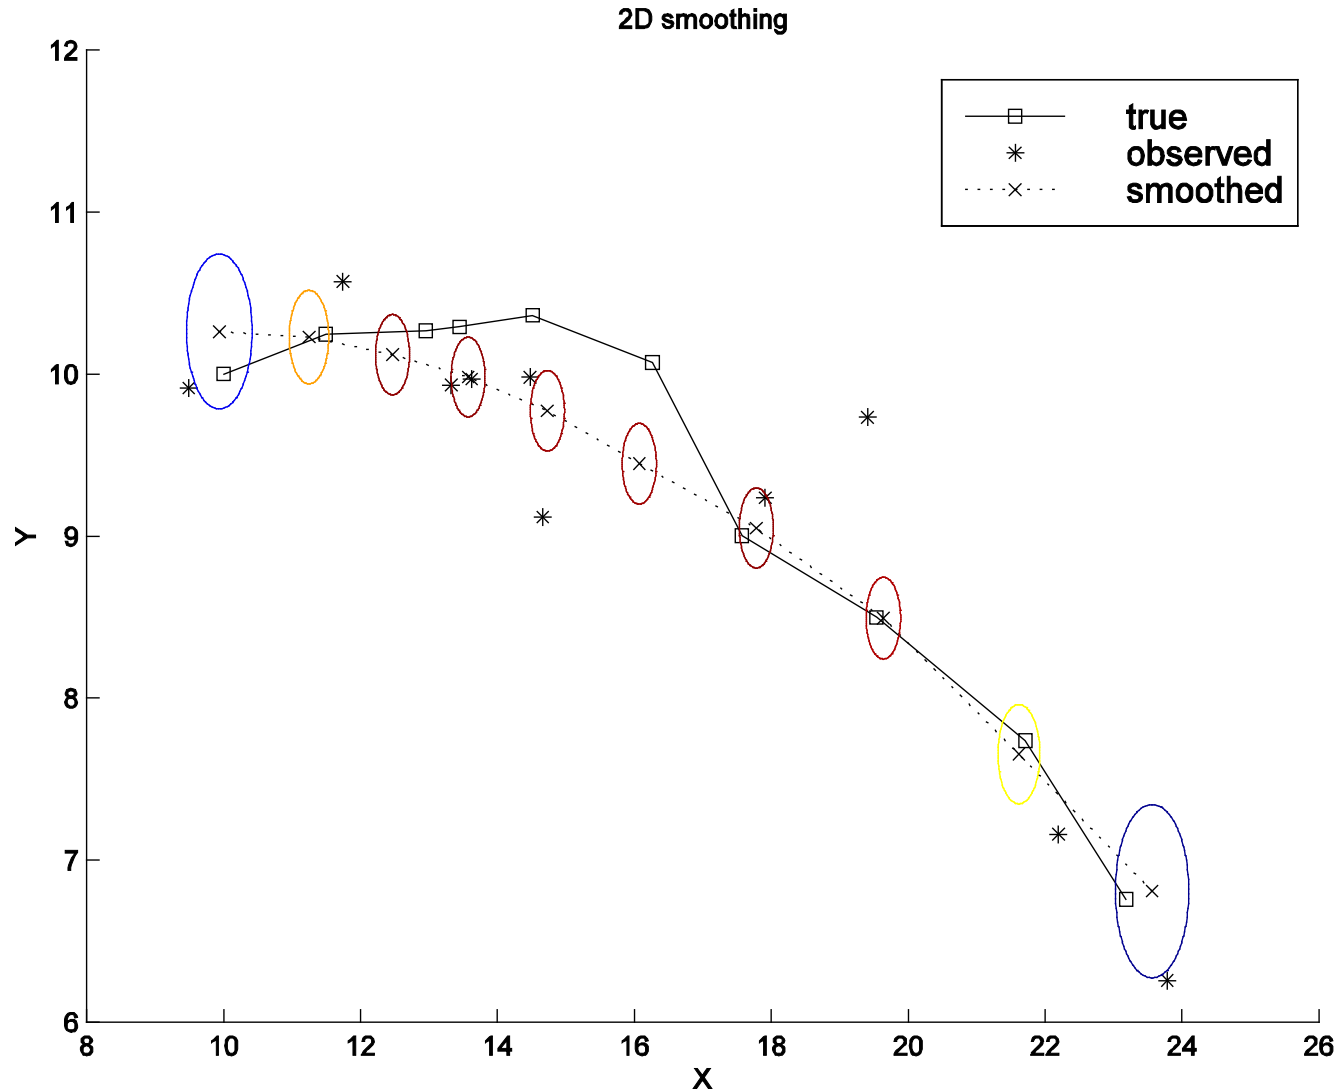
Σ_t and K_t are independent of the observed sequence and can thus be computed offline.

2D tracking: Filtering



$$\vec{X} = (X, dX/dt, Y, dY/dt)^T$$

2D tracking: Smoothing



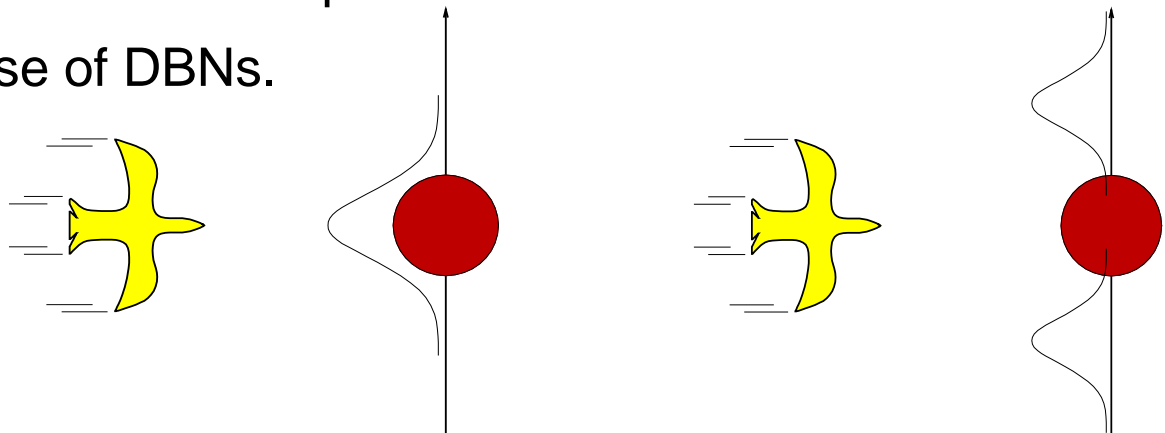
Simple Kalman filtering is not applicable if the transition model is non-linear.

Extension:

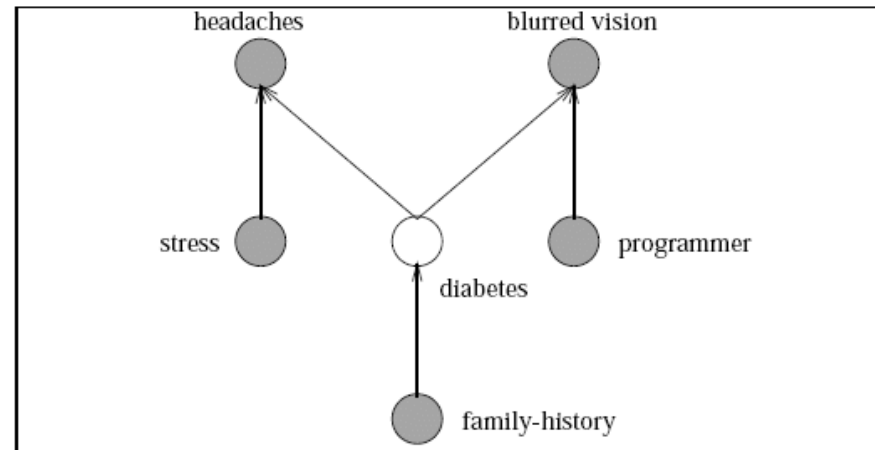
- Non-linearities can be treated by assuming **local linearity** in an environment of $x_t = \mu_t$.
- But this will fail if the system has a non-linearity at $x_t = \mu_t$.

Example: Bird is flying towards a tree.

- Solution: **Switching-Kalman-Filter**
 - Applies several filters in parallel
 - Special case of DBNs.



So far: Static Bayes networks for modeling dependencies without time:



[RN]

HMMs are a special case of **dynamic Bayesian networks (DBN)** with just **one variable**.

Kalman filters are a special case with **Gaussian** distributions.

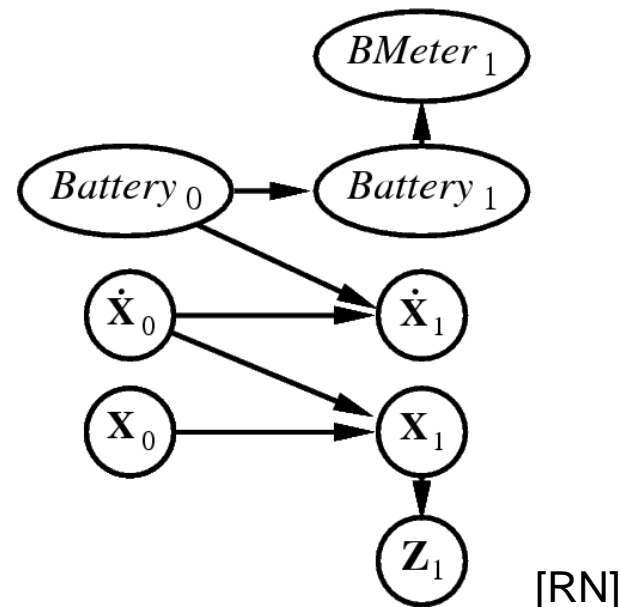
A **general DBN** is a temporal probability model with

- an arbitrary number of random variables X_t , and
- evidence variables E_t
- for each time step.

Example:

Robot with

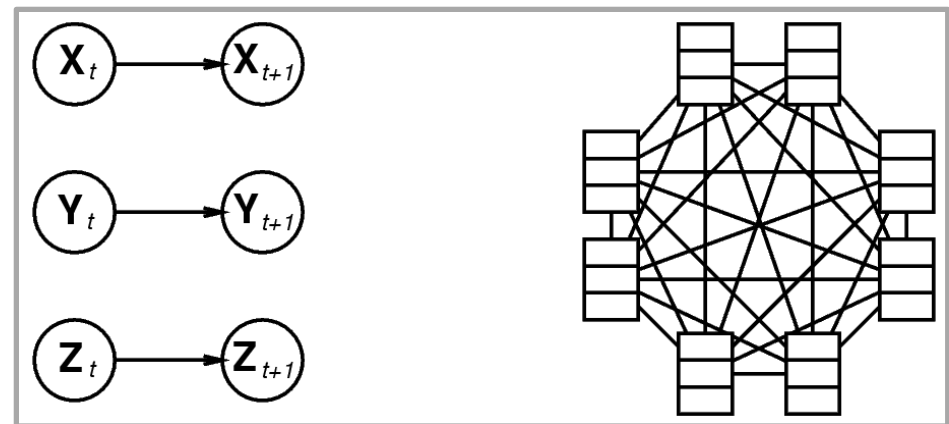
- state variables position X_t , speed V_t , and battery power $Battery_t$;
- evidence variables measured position Z_t and $BMeter_t$.
- for $t=0$ and $t=1$.



- Every HMM is a DBN with just one variable.
- Every DBN with discrete variables can be represented as a HMM:
 - Combine all variables of the DBN to a single HMM-Variable.
 - The HMM-variable has one value for each combination of the variables of the DBN.
 - Problem: Combinatorial explosion.
- DBN are much better suited than HMMs as they employ „factorized“ states with an exponentially smaller number of parameters.

Example: 20 boolean state variables with 3 parents each. Parameters:

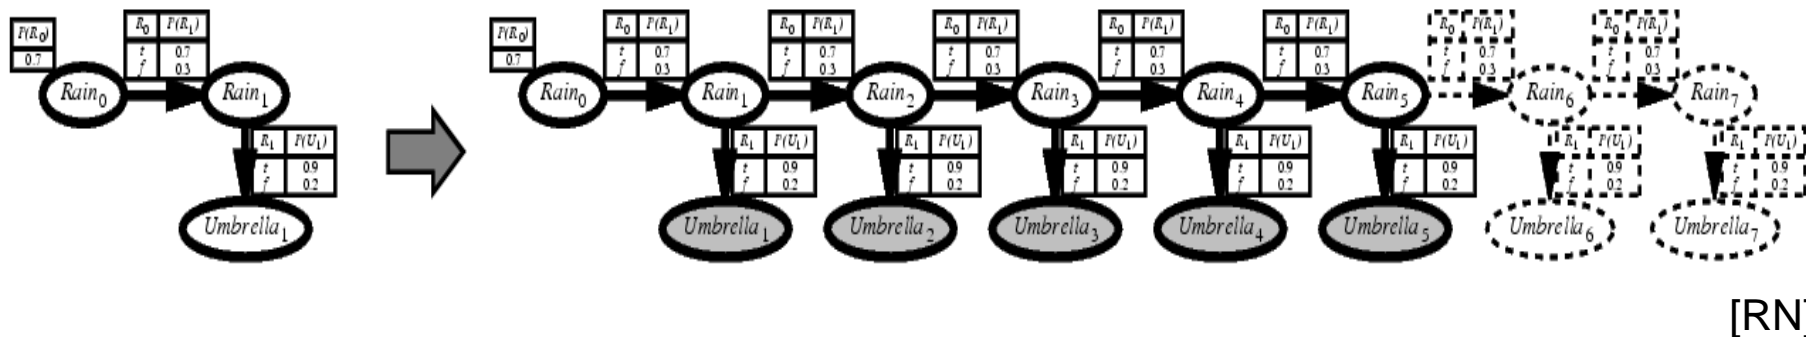
DBN $20 \times 2^3 = 160$,
HMM $2^{20} \times 2^{20}$.



Naive method:

- **Unroll** DBN (represent each time step explicitly).
- Apply Algorithm for static Bayes nets.

Problem: Memory and computational effort $O(t)$.



[RN]

Alternative: **Roll up filtering**

Add step $t+1$, then sum out the variables of step t .

Only possible (for realistic size) with approximation methods such as **particle filtering**.

- Temporal probability models represent a domain using random variables for hidden states and observable evidences.
- Three assumptions
 - Markov process,
 - Markov sensor,
 - Stationary domain.
- Several types of inference: Filtering, prediction, smoothing, most likely explanation (Viterbi algorithm).
- HMMs model a Markov process using a single variable.
- Kalman filtering employs an arbitrary number of state variables but only Gaussian distribution.
- DBNs have an arbitrary number of variables and arbitrary distributions, but exact inference is infeasible due to computational effort. Particle filtering is a good approximation for filtering.

Speech Recognition

- Speech recognition is an important application of temporal probability models.
- Recognize a sequence of words from a (raw) speech signal.
- Speech understanding:
 - Interpret sequence of words.
 - Find relation to other data, e.g., other sensors or a knowledge base.
- Speech signals are highly variable, ambiguous, noisy etc.
- Speech signals can not be classified after the simple scheme *signal* → *features* → *classifier* → *symbols*.
- Rather, simultaneous recognition on different levels of abstraction is required.

Task:

What is the **most likely word sequence** given a signal?

→ Choose *Words* such that $P(\text{Words} \mid \text{signal})$ is maximized.

Bayes rule:

$$P(\text{Words} \mid \text{signal}) = \alpha P(\text{signal} \mid \text{Words}) P(\text{Words}).$$

Thus the problem is decomposed into an **acoustic model** and a **language model**.

Words are the hidden state sequence, *signal* is the observation (evidence) sequence.

- For classification, a small number of different entities and large number of training samples of each is required.
- English has about 700000 words,
- consisting of 10000 syllables,
- but these consist of only 40-50 **phones** (speech sounds).

- **Phones** are formed by the **articulators** (lips, teeth, tongue, vocal cords, air flow).
- Phones are closer to the signal than words.
 - ➔ Acoustic model = pronunciation model + phone model
- **Phonemes** are the smallest units that have an effect on meaning (they do not carry meaning in isolation).
- Phonemes are combined to the smallest meaningful units: **Morphemes**.
- Phonemes \neq Characters
- **Allophones** are different speech sounds representing the same phoneme.
- Phonemes abstract phones to a representational level between signal and words.

DARPA-alphabet for American English (ARPAbet)

[iy]	b <u>e</u> at	[b]	<u>b</u> et	[p]	<u>p</u> et
[ih]	b <u>i</u> t	[ch]	<u>Ch</u> et	[r]	<u>r</u> at
[ey]	b <u>e</u> t	[d]	<u>d</u> ebt	[s]	<u>s</u> et
[ao]	b <u>o</u> ught	[hh]	<u>h</u> at	[th]	<u>th</u> ick
[ow]	b <u>o</u> at	[hv]	<u>h</u> igh	[dh]	<u>th</u> at
[er]	B <u>e</u> rt	[l]	<u>l</u> et	[w]	<u>w</u> et
[ix]	ros <u>e</u> s	[ng]	s <u>i</u> ng	[en]	butt <u>o</u> n
:	:	:	:	:	:

[RN]

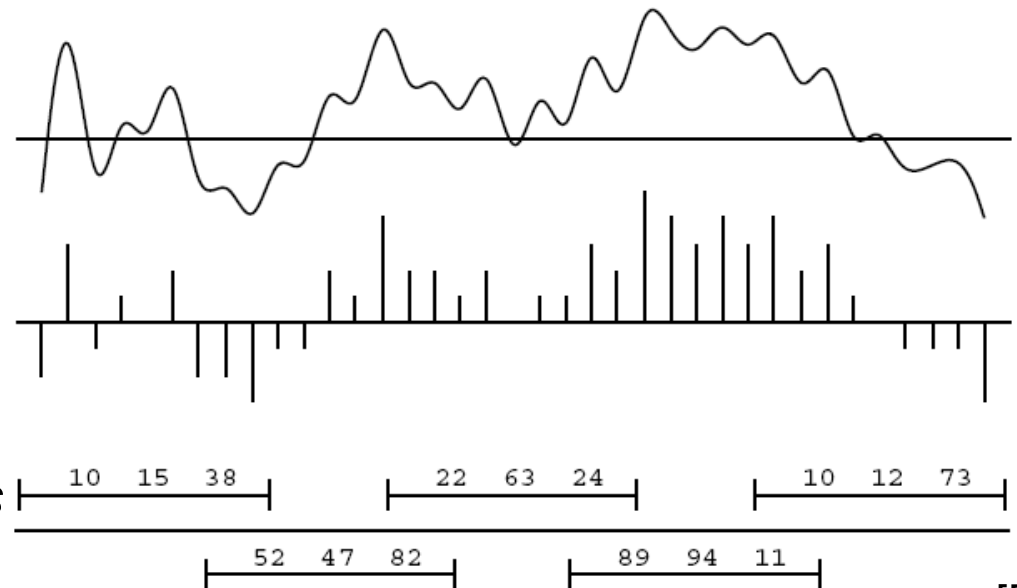
E.g. „ceiling“: [s iy l ich ng] / [s iy l ix ng] / [s iy l en]

- Signal: Displacement of microphone membrane as a function of time.
- Representation: 8-16 kHz sampling, 8-12 bit quantization.
- Signal is processed in overlapping **frames** of 30 ms.
- Data reduction: Each frame is represented by features.
- **Features**: E.g., peaks of the power spectrum.

Analog signal

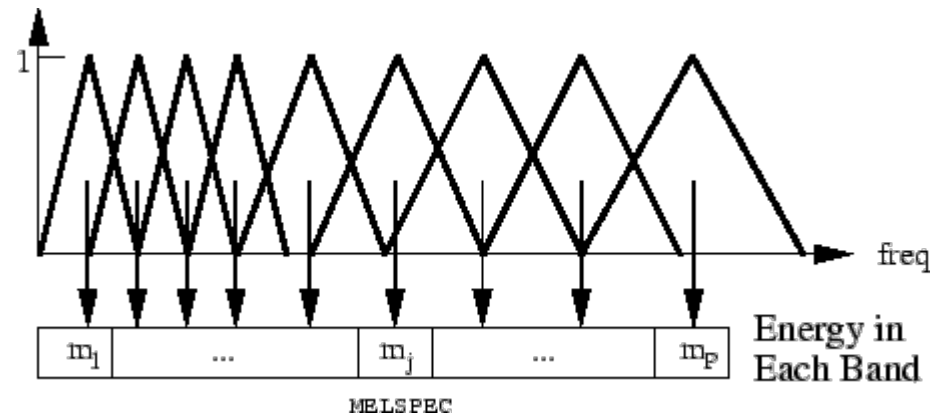
Sampled signal

Frames with features



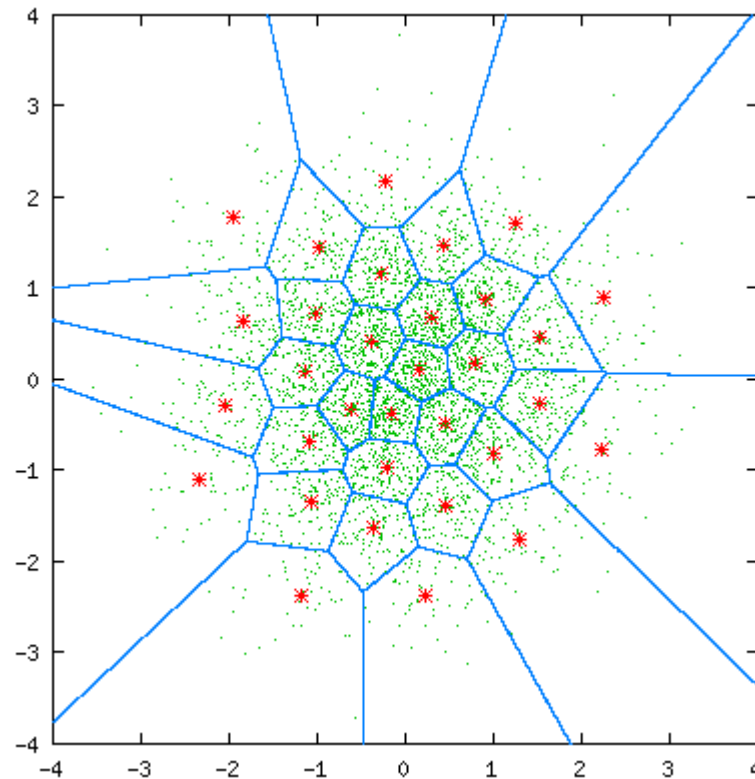
[RN]

- Overlap of frames 50%-75%.
- Features are, e.g., the distribution of energy over different frequencies, or change rates.
- Note energy distribution underlies uncertainty relation.
- Features may correlate with the activities of the articulators.

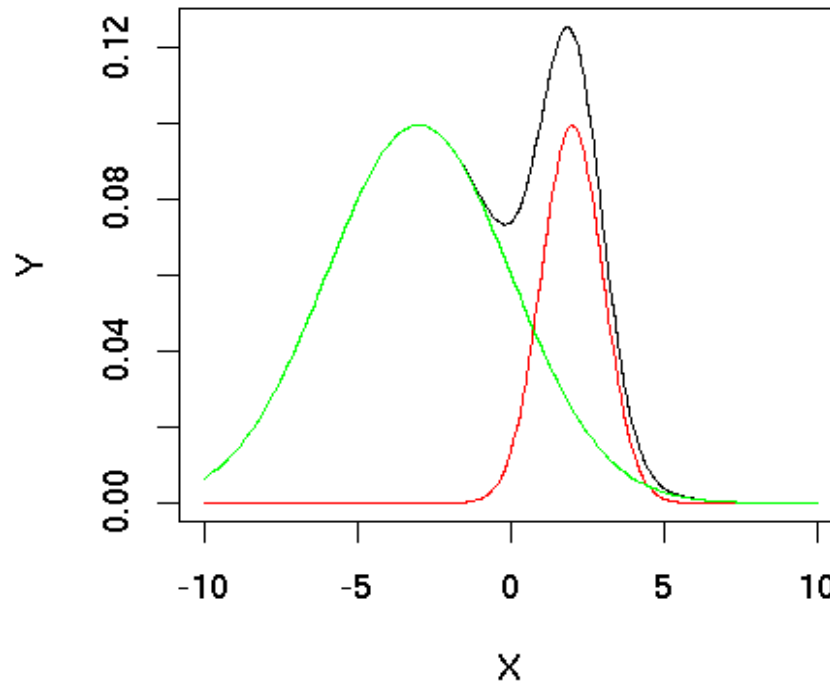


- Frame features are vectors of high dimensionality, leaving still many options to encode a phone.
- $P(\text{Features} \mid \text{Phone})$ represents the frame features.
- Better and more compact representation by, e.g.,
 - natural numbers obtained from clustering, or
 - parameters of a Gaussian mixture model

Clustering can be used to form groups of frequent features, natural numbers denote the groups (centers):



Gaussian mixtures describe $P(\text{Features} \mid \text{Phone})$ better than clusters.

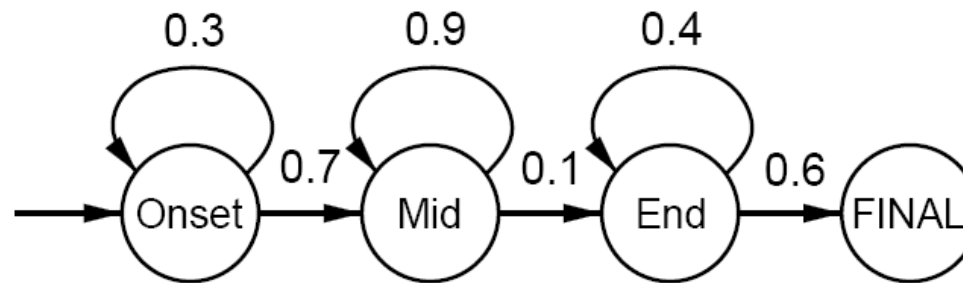


$$\frac{1}{4} N(2,1) + \frac{3}{4} N(-3,3)$$

- Phones exhibit inner structure.
- This structure can be modeled effectively a **three state phone model**:
 - Each phone consists of *Onset, Mid, End*.
 - Example: [t] has silent *Onset*, explosive *Mid*, hissing *End*.
 - Thus $P(\text{Features} \mid \text{Phone})$ is replaced by
 $P(\text{Features} \mid \text{Phone}, \text{Phase})$.

- Problem: Phones sound different, depending on neighboring speech sounds.
- These **coarticulation effects** come about because the articulators can not switch between positions instantaneously.
- Model for coarticulation: **Triphone context**
 - Each of n speech sounds is now represented by n^2 speech sounds which depend on both neighboring speech sounds.
 - Example: $[t]$ in „star“ is represented by $[t(s,aa)]$.
- Combining the three state model with the triphone model makes representation grow from n to n^3 , but this is worth the expense.

Phone HMM for [m]:

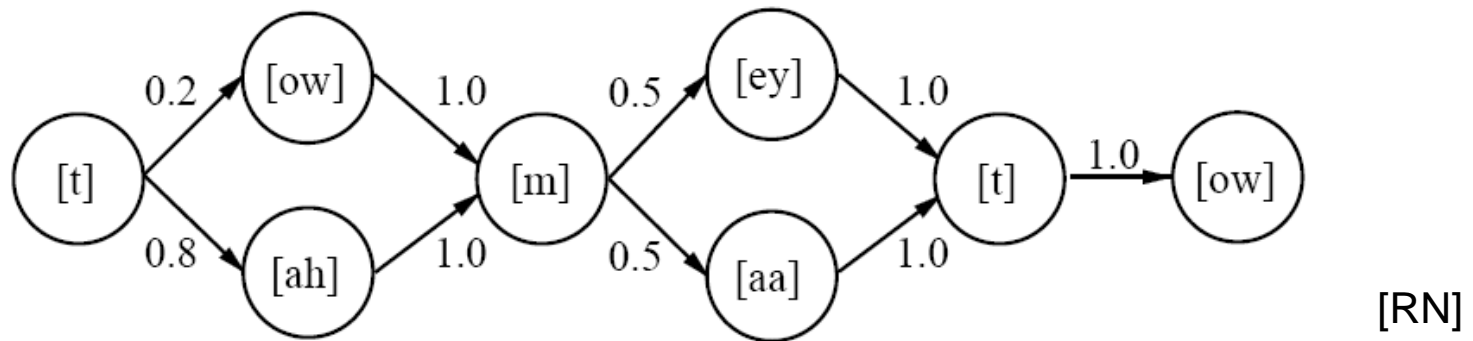


[RN]

To each of the three states of the **Phone HMM** belong output probabilities for the features (e.g., cluster numbers):

Onset:	Mid:	End:	
C1: 0.5	C3: 0.2	C4: 0.1	
C2: 0.2	C4: 0.7	C6: 0.5	
C3: 0.3	C5: 0.1	C7: 0.4	[RN]

A word is represented by a probability distribution over a phone sequence. This sequence is represented by a HMM:



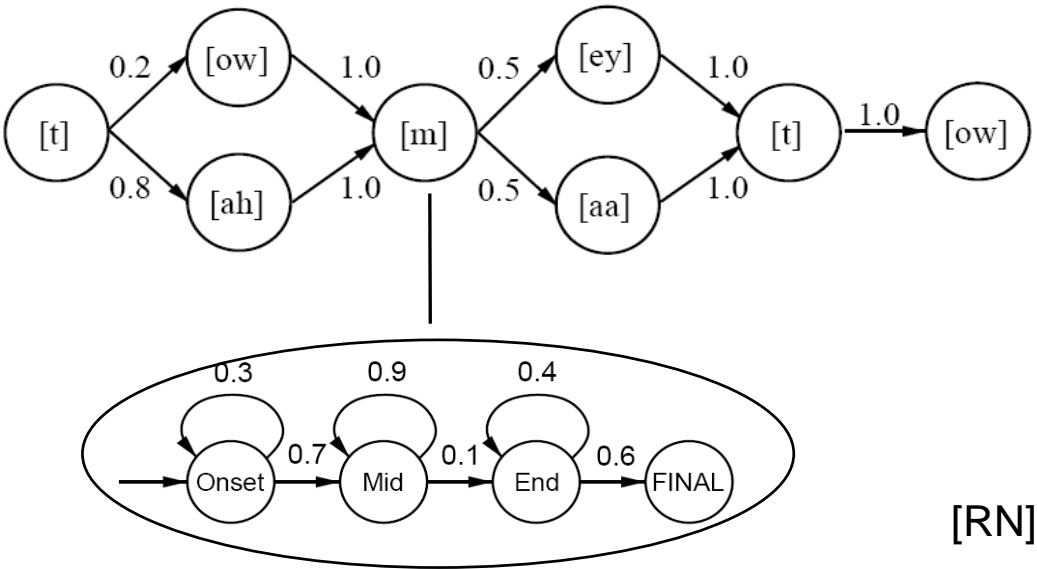
$$P([\text{towmeytow}] \mid \text{„tomato“}) = P([\text{towmaatow}] \mid \text{„tomato“}) = 0.1$$

$$P([\text{tahmeytow}] \mid \text{„tomato“}) = P([\text{tahmaatow}] \mid \text{„tomato“}) = 0.4$$

- The structure of the HMM is created manually.
- Transition probabilities are estimated from data.

A **word model** consists of the **phone models** and the **pronunciation model**.

Word model for *Tomato*:



State of a word HMM = phone + phone state,
e.g., the word HMM of *Tomato* has a state **[m]_{Mid}**.

Phone models + word models fix $P(e_{1:t} | \text{Word})$ for **isolated words**.
where $e_{1:t}$ are the observed features.

Recognizing a word means maximizing

$$P(\text{Word} | e_{1:t}) = \alpha P(e_{1:t} | \text{Word}) P(\text{Word}),$$

where the prior $P(\text{Word})$ is just obtained from the word frequencies.

$P(e_{1:t} | \text{Word})$ is computed recursively by

$$P(X_{t+1}, e_{1:t+1}) = \text{Forward}(P(X_t, e_{1:t}), e_{t+1})$$

and
$$P(e_{1:t} | \text{Word}) = \sum_{x_t} P(x_t, e_{1:t}).$$

Recognition of isolated words (e.g. for dictation) reaches 95-99% accuracy (with training on a particular person).

Recognition of continuous speech \neq recognition of sequence of isolated words, because

- adjacent words are strongly correlated,
- the most likely sequence of words \neq the sequence of most likely words,
- segmentation of words is difficult, because there are few gaps between words which become visible on the signal level (only on the high level of human speech processing),
- there is cross-word coarticulation, e.g., „next thing“.

Recognition of continuous speech manage 60-80% accuracy.

A language model specifies the a priori probability of each sequence of words using the chain rule:

$$P(w_1 \dots w_n) = \prod_{i=1}^n P(w_i \mid w_1 \dots w_{i-1}).$$

Most factors are hard to estimate.

Bigram model as an approximation:

$$P(w_i \mid w_1 \dots w_{i-1}) \approx P(w_i \mid w_{i-1}),$$

i.e., first order **Markov assumption**.

Training: Count all word pairs in a large text corpus.

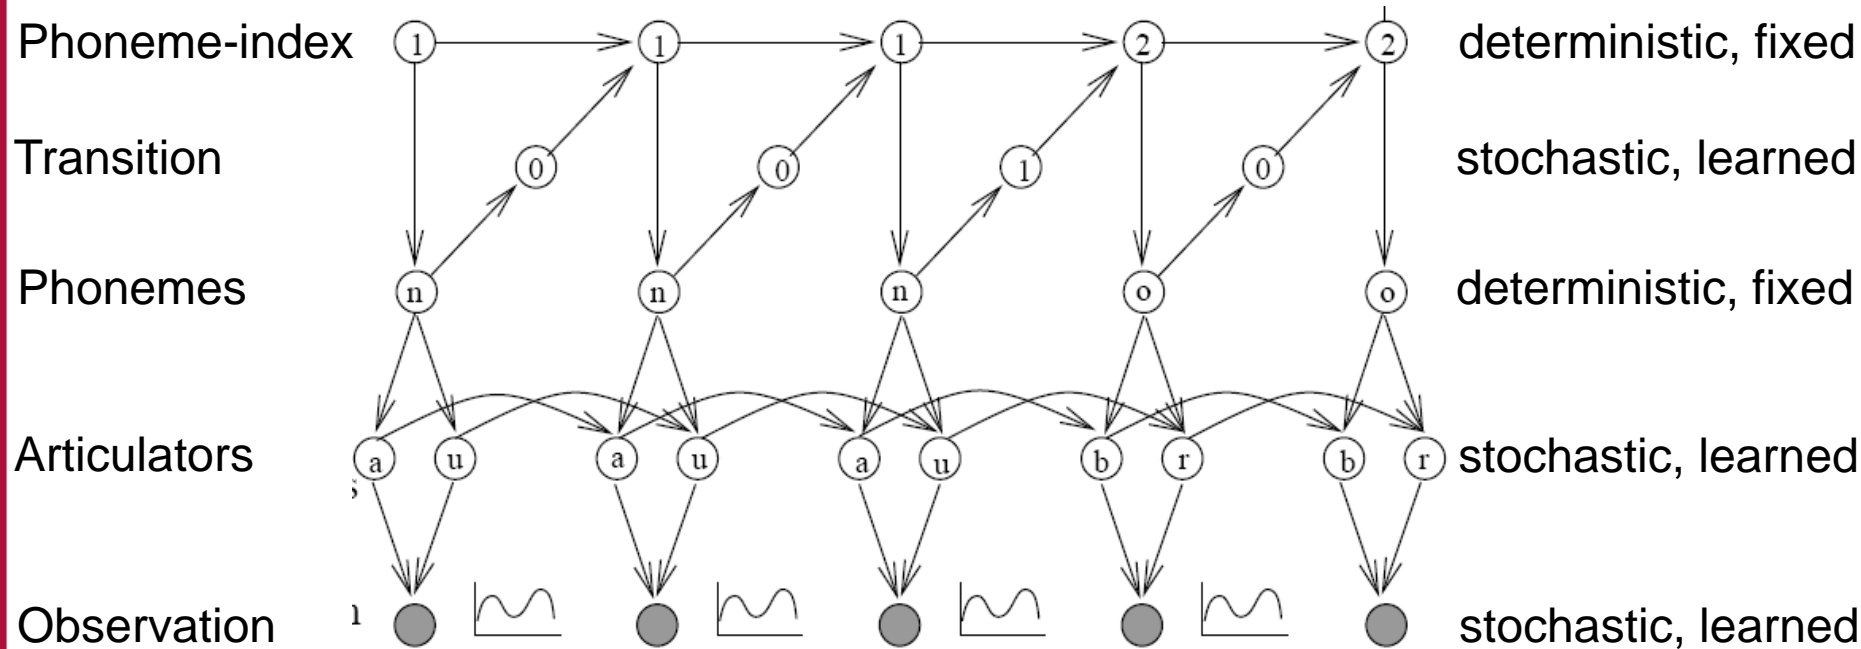
More complex models such as Trigrams

$$P(w_i \mid w_1 \dots w_{i-1}) \approx P(w_i \mid w_{i-1}, w_{i-2}),$$

or grammars lead to some improvement.

- Combine **word model** and Bigram **language model** to an HMM.
- States of the combined HMM are specified by **word**, **phone** and **phone state**.
- Example: $[m]^{\text{Tomato}}_{\text{Mid}}$.
- Transitions:
 - Phone state – phone state (within a phone),
 - Phone – phone (within a word),
 - Word final state – word initial state (between words).
- Representational effort:
The combined HMM for W words with an average of L three state phones has $3LW$ states.

- Most likely **phone sequence** is found by Viterbi algorithm
 - this fixes also the **word sequence** and
 - solves the **segmentation** problem.
- But: The word sequence obtained from the most likely phone sequence is not necessarily the most likely word sequence ...
- ... because probability of a word sequence = sum of probabilities of all corresponding state sequences.
- Solution: **A*-decoder** to find the most likely word sequence with moderate computational effort (Jelinek 1969).



- Further variables for gender, accent, speed are easy to add.
- Better performance than HMMs.

- Speech recognition has been formulated as probabilistic inference since 70ies.
- Evidence = speech signal
- Hidden variables = phone and word sequences
- Context effects such as coarticulation are handled by augmenting the states.
- Highly successful approach.

- [M] Online material available at www.cs.cmu.edu/~tom/mlbook.html for the textbook: Tom M. Mitchell: *Machine Learning*, McGraw-Hill
- [RN] Stuart Russell, Peter Norvig: *Artificial Intelligence*, Pearson
- [H] Gunther Heidemann, 2012.

- p-norm unit circles
- Optimization based clustering
- K-Means
- Conceptual clustering
- Hebbian Learning:
 - Hebb rule
 - Anti-Hebb rule
- Eigenfaces
- Principal curves and SOM

- MLP
 - Parameters
 - Comparison to RBF
- RBF
 - Parameters
- SOM
- Q-Learning: Probabilistic choice of actions